

Collaboration and Competition

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

In [3]:

```
!pip -q install ./python
```

The environment is already saved in the Workspace and can be accessed at the file path provided below.

In [4]:

```
from unityagents import UnityEnvironment
# You are welcome to use this coding environment to train your agent for the project.
# Follow the instructions below to get started!

import numpy as np

env = UnityEnvironment(file_name="/data/Tennis_Linux_NoVis/Tennis")
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
```

```
Unity brain name: TennisBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 8
    Number of stacked Vector Observation: 3
    Vector Action space type: continuous
    Vector Action space size (per agent): 2
    Vector Action descriptions: ,
```

Environments contain **brains** which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

In [5]:

```
# get the default brain
brain_name = env.brain_names[0]
brain = env.brains[brain_name]
```

2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

In [6]:

```
# reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents
num_agents = len(env_info.agents)
print('Number of agents:', num_agents)

# size of each action
action_size = brain.vector_action_space_size
print('Size of each action:', action_size)

# examine the state space
states = env_info.vector_observations
state_size = states.shape[1]
print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
print('The state for the first agent looks like:', states[0])
```

```
Number of agents: 2
Size of each action: 2
There are 2 agents. Each observes a state with length: 24
The state for the first agent looks like: [ 0.          0.          0.
 0.          0.          0.          0.
 0.          0.          0.          0.          0.          0.
 0.          0.          -6.65278625 -1.5          -0.          0.
 6.83172083  6.          -0.          0.          ]
```

3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agents while they are training**, and you should set `train_mode=True` to restart the environment.

In [7]:

```

for i in range(5):                                # play game for 5 episodes
    env_info = env.reset(train_mode=False)[brain_name] # reset the environment
    states = env_info.vector_observations            # get the current state (for
each agent)
    scores = np.zeros(num_agents)                  # initialize the score (for
each agent)
    while True:
        actions = np.random.randn(num_agents, action_size) # select an action (for each
agent)
        actions = np.clip(actions, -1, 1)           # all actions between -1 and
1
        env_info = env.step(actions)[brain_name]    # send all actions to the en
vironment
        next_states = env_info.vector_observations  # get next state (for each a
gent)
        rewards = env_info.rewards                 # get reward (for each agen
t)
        dones = env_info.local_done                 # see if episode finished
        scores += env_info.rewards                  # update the score (for each
agent)
        states = next_states                        # roll over states to next t
ime step
        if np.any(dones):                          # exit loop if episode finis
hed
            break
    print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores
)))

```

```

Total score (averaged over agents) this episode: -0.004999999888241291
Total score (averaged over agents) this episode: -0.004999999888241291
Total score (averaged over agents) this episode: -0.004999999888241291
Total score (averaged over agents) this episode: -0.004999999888241291
Total score (averaged over agents) this episode: -0.004999999888241291

```

When finished, you can close the environment.

In []:

```
env.close()
```

4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**:

- When training the environment, set `train_mode=True` , so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on ***Jupyter*** in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agents while they are training. However, ***after training the agents***, you can download the saved model weights to watch the agents on your own machine!

In [8]:

```
import numpy as np

import torch
import torch.nn as nn
import torch.nn.functional as F

def hidden_init(layer):
    fan_in = layer.weight.data.size()[0]
    lim = 1. / np.sqrt(fan_in)
    return (-lim, lim)

class Actor(nn.Module):
    def __init__(self, state_size, action_size, seed, fc1_units=256, fc2_units=128):
        super(Actor, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fc1 = nn.Linear(state_size*2, fc1_units)
        self.fc2 = nn.Linear(fc1_units, fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)
        self.reset_parameters()

    def reset_parameters(self):
        self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state):
        x = F.relu(self.fc1(state))
        x = F.relu(self.fc2(x))
        return torch.tanh(self.fc3(x))

class Critic(nn.Module):
    def __init__(self, state_size, action_size, seed, fcs1_units=256, fc2_units=128):
        super(Critic, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fcs1 = nn.Linear(state_size*2, fcs1_units)
        self.fc2 = nn.Linear(fcs1_units+(action_size*2), fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)
        self.reset_parameters()

    def reset_parameters(self):
        self.fcs1.weight.data.uniform_(*hidden_init(self.fcs1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state, action):
        xs = F.relu(self.fcs1(state))
        x = torch.cat((xs, action), dim=1)
        x = F.relu(self.fc2(x))
        return self.fc3(x)
```

In [9]:

```

import numpy as np
import random
import copy
from collections import namedtuple, deque

import torch
import torch.nn.functional as F
import torch.optim as optim

BUFFER_SIZE = int(1e6)  # replay buffer size
BATCH_SIZE = 128        # minibatch size
LR_ACTOR = 1e-3         # learning rate of the actor
LR_CRITIC = 1e-3        # learning rate of the critic
WEIGHT_DECAY = 0        # L2 weight decay
LEARN_EVERY = 1         # learning timestep interval
LEARN_NUM = 5           # number of learning passes
GAMMA = 0.99            # discount factor
TAU = 8e-3              # for soft update of target parameters
OU_SIGMA = 0.2          # Ornstein-Uhlenbeck noise parameter, volatility
OU_THETA = 0.15         # Ornstein-Uhlenbeck noise parameter, speed of mean reversion
EPS_START = 5.0         # initial value for epsilon in noise decay process in Agent.act
()
EPS_EP_END = 300        # episode to end the noise decay process
EPS_FINAL = 0           # final value for epsilon after decay

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class Agent():
    def __init__(self, state_size, action_size, num_agents, random_seed):
        self.state_size = state_size
        self.action_size = action_size
        self.num_agents = num_agents
        self.seed = random.seed(random_seed)
        self.eps = EPS_START
        self.eps_decay = 1/(EPS_EP_END*LEARN_NUM)  # set decay rate based on epsilon en
d target
        self.timestep = 0

        # Actor Network (w/ Target Network)
        self.actor_local = Actor(state_size, action_size, random_seed).to(device)
        self.actor_target = Actor(state_size, action_size, random_seed).to(device)
        self.actor_optimizer = optim.Adam(self.actor_local.parameters(), lr=LR_ACTOR)

        # Critic Network (w/ Target Network)
        self.critic_local = Critic(state_size, action_size, random_seed).to(device)
        self.critic_target = Critic(state_size, action_size, random_seed).to(device)
        self.critic_optimizer = optim.Adam(self.critic_local.parameters(), lr=LR_CRITIC
, weight_decay=WEIGHT_DECAY)

        # Noise process
        self.noise = OUNoise((num_agents, action_size), random_seed)

        # Replay memory
        self.memory = ReplayBuffer(action_size, BUFFER_SIZE, BATCH_SIZE, random_seed)

    def step(self, state, action, reward, next_state, done, agent_number):
        self.timestep += 1

```

```

# Save experience / reward
self.memory.add(state, action, reward, next_state, done)
# Learn, if enough samples are available in memory and at learning interval set
tings
if len(self.memory) > BATCH_SIZE and self.timestep % LEARN_EVERY == 0:
    for _ in range(LEARN_NUM):
        experiences = self.memory.sample()
        self.learn(experiences, GAMMA, agent_number)

def act(self, states, add_noise):
    states = torch.from_numpy(states).float().to(device)
    actions = np.zeros((self.num_agents, self.action_size))
    self.actor_local.eval()
    with torch.no_grad():
        # get action for each agent and concatenate them
        for agent_num, state in enumerate(states):
            action = self.actor_local(state).cpu().data.numpy()
            actions[agent_num, :] = action
    self.actor_local.train()
    # add noise to actions
    if add_noise:
        actions += self.eps * self.noise.sample()
    actions = np.clip(actions, -1, 1)
    return actions

def reset(self):
    self.noise.reset()

def learn(self, experiences, gamma, agent_number):
    states, actions, rewards, next_states, dones = experiences

    # ----- update critic ----- #
    # Get predicted next-state actions and Q values from target models
    actions_next = self.actor_target(next_states)
    # Construct next actions vector relative to the agent
    if agent_number == 0:
        actions_next = torch.cat((actions_next, actions[:,2:]), dim=1)
    else:
        actions_next = torch.cat((actions[:,2:], actions_next), dim=1)
    # Compute Q targets for current states (y_i)
    Q_targets_next = self.critic_target(next_states, actions_next)
    Q_targets = rewards + (gamma * Q_targets_next * (1 - dones))
    # Compute critic loss
    Q_expected = self.critic_local(states, actions)
    critic_loss = F.mse_loss(Q_expected, Q_targets)
    # Minimize the loss
    self.critic_optimizer.zero_grad()
    critic_loss.backward()
    torch.nn.utils.clip_grad_norm_(self.critic_local.parameters(), 1)
    self.critic_optimizer.step()

    # ----- update actor ----- #
    # Compute actor loss
    actions_pred = self.actor_local(states)
    # Construct action prediction vector relative to each agent
    if agent_number == 0:
        actions_pred = torch.cat((actions_pred, actions[:,2:]), dim=1)
    else:
        actions_pred = torch.cat((actions[:,2:], actions_pred), dim=1)
    # Compute actor loss
    actor_loss = -self.critic_local(states, actions_pred).mean()

```

```

    # Minimize the Loss
    self.actor_optimizer.zero_grad()
    actor_loss.backward()
    self.actor_optimizer.step()

    # ----- update target networks ----- #
    self.soft_update(self.critic_local, self.critic_target, TAU)
    self.soft_update(self.actor_local, self.actor_target, TAU)

    # update noise decay parameter
    self.eps -= self.eps_decay
    self.eps = max(self.eps, EPS_FINAL)
    self.noise.reset()

def soft_update(self, local_model, target_model, tau):
    for target_param, local_param in zip(target_model.parameters(), local_model.parameters()):
        target_param.data.copy_(tau*local_param.data + (1.0-tau)*target_param.data)

class OUNoise:

    def __init__(self, size, seed, mu=0.0, theta=OU_THETA, sigma=OU_SIGMA):
        self.mu = mu * np.ones(size)
        self.theta = theta
        self.sigma = sigma
        self.seed = random.seed(seed)
        self.size = size
        self.reset()

    def reset(self):
        self.state = copy.copy(self.mu)

    def sample(self):
        x = self.state
        dx = self.theta * (self.mu - x) + self.sigma * np.random.standard_normal(self.size)
        self.state = x + dx
        return self.state

class ReplayBuffer:

    def __init__(self, action_size, buffer_size, batch_size, seed):
        self.action_size = action_size
        self.memory = deque(maxlen=buffer_size) # internal memory (deque)
        self.batch_size = batch_size
        self.experience = namedtuple("Experience", field_names=["state", "action", "reward", "next_state", "done"])
        self.seed = random.seed(seed)

    def add(self, state, action, reward, next_state, done):
        """Add a new experience to memory."""
        e = self.experience(state, action, reward, next_state, done)
        self.memory.append(e)

    def sample(self):
        experiences = random.sample(self.memory, k=self.batch_size)

        states = torch.from_numpy(np.vstack([e.state for e in experiences if e is not None])).float().to(device)
        actions = torch.from_numpy(np.vstack([e.action for e in experiences if e is not None])).float().to(device)
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences if e is not None])).float().to(device)

```



```
None])).float().to(device)
    next_states = torch.from_numpy(np.vstack([e.next_state for e in experiences if
e is not None])).float().to(device)
    dones = torch.from_numpy(np.vstack([e.done for e in experiences if e is not None])).astype(np.uint8).float().to(device)

    return (states, actions, rewards, next_states, dones)

def __len__(self):
    return len(self.memory)
```

In [10]:

```

SOLVED_SCORE = 2.7
CONSEC_EPISODES = 100
PRINT_EVERY = 10
ADD_NOISE = True

def maddpg(n_episodes=6000, max_t=1000, train_mode=True):
    scores_window = deque(maxlen=CONSEC_EPISODES)
    scores_all = []
    moving_average = []
    best_score = -np.inf
    best_episode = 0
    already_solved = False

    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=train_mode)[brain_name]           # reset the env
        states = np.reshape(env_info.vector_observations, (1,48)) # get states and combine them
        agent_0.reset()
        agent_1.reset()
        scores = np.zeros(num_agents)
        while True:
            actions = get_actions(states, ADD_NOISE) # choose agent actions and combine them
            env_info = env.step(actions)[brain_name] # send both agents' actions together to the environment
            next_states = np.reshape(env_info.vector_observations, (1, 48)) # combine the agent next states
            rewards = env_info.rewards # get reward
            done = env_info.local_done # see if episode finished
            agent_0.step(states, actions, rewards[0], next_states, done, 0) # agent 1 learns
            agent_1.step(states, actions, rewards[1], next_states, done, 1) # agent 2 learns
            scores += np.max(rewards) # update the score for each agent
            states = next_states # roll over states to next time step
            if np.any(done): # exit loop if episode finished
                break

            ep_best_score = np.max(scores)
            scores_window.append(ep_best_score)
            scores_all.append(ep_best_score)
            moving_average.append(np.mean(scores_window))

            # save best score
            if ep_best_score > best_score:
                best_score = ep_best_score
                best_episode = i_episode

        # if (i_episode > 1000 and np.mean(scores_window) < 0.08):
        #     break

    # print results
    if i_episode % 10 == 0:

```

```

        print('\rEpisodes {:0>4d}-{:0>4d}\tMax Reward: {:.3f}\tMoving Average: {:.3f}'.format(
            i_episode-PRINT_EVERY, i_episode, np.max(scores_all[-PRINT_EVERY:]), moving_average[-1]))
#         torch.save(agent_0.actor_local.state_dict(), 'checkpoint_actor_0.pth')
#         torch.save(agent_0.critic_local.state_dict(), 'checkpoint_critic_0.pth')
#         torch.save(agent_1.actor_local.state_dict(), 'checkpoint_actor_1.pth')
#         torch.save(agent_1.critic_local.state_dict(), 'checkpoint_critic_1.pth')

    if i_episode % 100 == 0:
        print('\r\nEpisodes {:0>4d}-{:0>4d}\tMax Reward: {:.3f}\tMoving Average: {:.3f} \n Saved!!!\n'.format(
            i_episode-PRINT_EVERY, i_episode, np.max(scores_all[-PRINT_EVERY:]), moving_average[-1]))
        torch.save(agent_0.actor_local.state_dict(), 'checkpoint_actor_0.pth')
        torch.save(agent_0.critic_local.state_dict(), 'checkpoint_critic_0.pth')
        torch.save(agent_1.actor_local.state_dict(), 'checkpoint_actor_1.pth')
        torch.save(agent_1.critic_local.state_dict(), 'checkpoint_critic_1.pth')

# determine if environment is solved and keep best performing models
    if moving_average[-1] >= 2.5:
        if not already_solved:
            print('\r<-- Environment solved in {:d} episodes! \n<-- Moving Average: {:.3f} over past {:d} episodes'.format(
                i_episode-CONSEC_EPISODES, moving_average[-1], CONSEC_EPISODES))
            already_solved = True
            # save weights
            torch.save(agent_0.actor_local.state_dict(), 'checkpoint_actor_0.pth')
            torch.save(agent_0.critic_local.state_dict(), 'checkpoint_critic_0.pth')
        )

        torch.save(agent_1.actor_local.state_dict(), 'checkpoint_actor_1.pth')
        torch.save(agent_1.critic_local.state_dict(), 'checkpoint_critic_1.pth')
    )

    elif ep_best_score >= best_score:
        print('\r<-- Best episode so far!\nEpisode {:0>4d}\tMax Reward: {:.3f}\tMoving Average: {:.3f}'.format(
            i_episode, ep_best_score, moving_average[-1]))
        # save weights
        torch.save(agent_0.actor_local.state_dict(), 'checkpoint_actor_0.pth')
        torch.save(agent_0.critic_local.state_dict(), 'checkpoint_critic_0.pth')
    )

        torch.save(agent_1.actor_local.state_dict(), 'checkpoint_actor_1.pth')
        torch.save(agent_1.critic_local.state_dict(), 'checkpoint_critic_1.pth')
    )

    elif (moving_average[-1]) >= 2.5:
        # stop training if model stops converging
        print('<-- Training stopped. Best score not matched or exceeded for 200 episodes')
        break
    else:
        continue

    return scores_all, moving_average

def get_actions(states, add_noise):
    '''gets actions for each agent and then combines them into one array'''
    action_0 = agent_0.act(states, add_noise) # agent 0 chooses an action
    action_1 = agent_1.act(states, add_noise) # agent 1 chooses an action
    return np.concatenate((action_0, action_1), axis=0).flatten()

```

In [11]:

```
# initialize agents
agent_0 = Agent(state_size, action_size, num_agents=1, random_seed=0)
agent_1 = Agent(state_size, action_size, num_agents=1, random_seed=0)

# Load the weights from file
# agent_0_weights = 'checkpoint_actor_0.pth'
# agent_1_weights = 'checkpoint_actor_1.pth'
# agent_0.actor_local.load_state_dict(torch.load(agent_0_weights))
# agent_1.actor_local.load_state_dict(torch.load(agent_1_weights))

# run the training loop

BUFFER_SIZE = int(1e6) # replay buffer size
BATCH_SIZE = 128      # minibatch size
LR_ACTOR = 1e-4       # learning rate of the actor
LR_CRITIC = 2e-4      # learning rate of the critic
WEIGHT_DECAY = 0      # L2 weight decay
LEARN_EVERY = 1       # learning timestep interval
LEARN_NUM = 1         # number of learning passes
GAMMA = 0.99          # discount factor
TAU = 7e-2            # for soft update of target parameters
OU_SIGMA = 0.2         # Ornstein-Uhlenbeck noise parameter, volatility
OU_THETA = 0.12        # Ornstein-Uhlenbeck noise parameter, speed of mean reversion
EPS_START = 5.5        # initial value for epsilon in noise decay process in A
gent.act()
EPS_EP_END = 250       # episode to end the noise decay process
EPS_FINAL = 0          # final value for epsilon after decay

scores, avgs = maddpg()
```

Episodes 0000-0010	Max Reward: 0.500	Moving Average: 0.060
Episodes 0010-0020	Max Reward: 0.000	Moving Average: 0.030
Episodes 0020-0030	Max Reward: 0.000	Moving Average: 0.020
Episodes 0030-0040	Max Reward: 0.000	Moving Average: 0.015
Episodes 0040-0050	Max Reward: 0.000	Moving Average: 0.012
Episodes 0050-0060	Max Reward: 0.000	Moving Average: 0.010
Episodes 0060-0070	Max Reward: 0.000	Moving Average: 0.009
Episodes 0070-0080	Max Reward: 0.000	Moving Average: 0.008
Episodes 0080-0090	Max Reward: 0.000	Moving Average: 0.007
Episodes 0090-0100	Max Reward: 0.100	Moving Average: 0.007

Episodes 0090-0100 Max Reward: 0.100 Moving Average: 0.007
Saved!!!

Episodes 0100-0110	Max Reward: 0.000	Moving Average: 0.001
Episodes 0110-0120	Max Reward: 0.000	Moving Average: 0.001
Episodes 0120-0130	Max Reward: 0.100	Moving Average: 0.002
Episodes 0130-0140	Max Reward: 0.100	Moving Average: 0.003
Episodes 0140-0150	Max Reward: 0.100	Moving Average: 0.006
Episodes 0150-0160	Max Reward: 0.300	Moving Average: 0.012
Episodes 0160-0170	Max Reward: 0.100	Moving Average: 0.014
Episodes 0170-0180	Max Reward: 0.300	Moving Average: 0.017
Episodes 0180-0190	Max Reward: 0.100	Moving Average: 0.019
Episodes 0190-0200	Max Reward: 0.100	Moving Average: 0.020

Episodes 0190-0200 Max Reward: 0.100 Moving Average: 0.020
Saved!!!

Episodes 0200-0210	Max Reward: 0.300	Moving Average: 0.027
Episodes 0210-0220	Max Reward: 0.400	Moving Average: 0.034
Episodes 0220-0230	Max Reward: 0.100	Moving Average: 0.035
Episodes 0230-0240	Max Reward: 0.100	Moving Average: 0.035
Episodes 0240-0250	Max Reward: 0.100	Moving Average: 0.034
Episodes 0250-0260	Max Reward: 0.100	Moving Average: 0.029
Episodes 0260-0270	Max Reward: 0.400	Moving Average: 0.039
Episodes 0270-0280	Max Reward: 0.200	Moving Average: 0.042
Episodes 0280-0290	Max Reward: 0.300	Moving Average: 0.047
Episodes 0290-0300	Max Reward: 0.300	Moving Average: 0.057

Episodes 0290-0300 Max Reward: 0.300 Moving Average: 0.057
Saved!!!

Episodes 0300-0310	Max Reward: 0.600	Moving Average: 0.069
Episodes 0310-0320	Max Reward: 0.400	Moving Average: 0.076
Episodes 0320-0330	Max Reward: 0.400	Moving Average: 0.093
Episodes 0330-0340	Max Reward: 0.300	Moving Average: 0.100
Episodes 0340-0350	Max Reward: 0.400	Moving Average: 0.111
Episodes 0350-0360	Max Reward: 0.500	Moving Average: 0.123
Episodes 0360-0370	Max Reward: 0.400	Moving Average: 0.123
Episodes 0370-0380	Max Reward: 0.400	Moving Average: 0.131
Episodes 0380-0390	Max Reward: 0.200	Moving Average: 0.131
Episodes 0390-0400	Max Reward: 0.900	Moving Average: 0.145

Episodes 0390-0400 Max Reward: 0.900 Moving Average: 0.145
Saved!!!

Episodes 0400-0410	Max Reward: 0.400	Moving Average: 0.142
Episodes 0410-0420	Max Reward: 0.200	Moving Average: 0.137
Episodes 0420-0430	Max Reward: 0.300	Moving Average: 0.130
Episodes 0430-0440	Max Reward: 0.400	Moving Average: 0.138
Episodes 0440-0450	Max Reward: 0.300	Moving Average: 0.140

Episodes 0450-0460	Max Reward: 0.400	Moving Average: 0.148
Episodes 0460-0470	Max Reward: 0.400	Moving Average: 0.148
Episodes 0470-0480	Max Reward: 0.400	Moving Average: 0.156
Episodes 0480-0490	Max Reward: 0.400	Moving Average: 0.161
Episodes 0490-0500	Max Reward: 0.400	Moving Average: 0.155
Episodes 0490-0500 Saved!!!	Max Reward: 0.400	Moving Average: 0.155
Episodes 0500-0510	Max Reward: 0.200	Moving Average: 0.149
Episodes 0510-0520	Max Reward: 0.400	Moving Average: 0.155
Episodes 0520-0530	Max Reward: 0.400	Moving Average: 0.155
Episodes 0530-0540	Max Reward: 0.200	Moving Average: 0.150
Episodes 0540-0550	Max Reward: 0.400	Moving Average: 0.146
Episodes 0550-0560	Max Reward: 0.400	Moving Average: 0.137
Episodes 0560-0570	Max Reward: 0.400	Moving Average: 0.149
Episodes 0570-0580	Max Reward: 0.400	Moving Average: 0.150
Episodes 0580-0590	Max Reward: 0.100	Moving Average: 0.142
Episodes 0590-0600	Max Reward: 0.400	Moving Average: 0.140
Episodes 0590-0600 Saved!!!	Max Reward: 0.400	Moving Average: 0.140
Episodes 0600-0610	Max Reward: 0.400	Moving Average: 0.149
Episodes 0610-0620	Max Reward: 0.400	Moving Average: 0.152
Episodes 0620-0630	Max Reward: 0.400	Moving Average: 0.159
Episodes 0630-0640	Max Reward: 0.400	Moving Average: 0.170
Episodes 0640-0650	Max Reward: 0.400	Moving Average: 0.183
Episodes 0650-0660	Max Reward: 0.400	Moving Average: 0.198
Episodes 0660-0670	Max Reward: 0.400	Moving Average: 0.199
Episodes 0670-0680	Max Reward: 0.400	Moving Average: 0.203
Episodes 0680-0690	Max Reward: 0.400	Moving Average: 0.221
Episodes 0690-0700	Max Reward: 0.400	Moving Average: 0.224
Episodes 0690-0700 Saved!!!	Max Reward: 0.400	Moving Average: 0.224
Episodes 0700-0710	Max Reward: 0.400	Moving Average: 0.228
Episodes 0710-0720	Max Reward: 0.400	Moving Average: 0.236
Episodes 0720-0730	Max Reward: 0.400	Moving Average: 0.238
Episodes 0730-0740	Max Reward: 0.400	Moving Average: 0.245
Episodes 0740-0750	Max Reward: 0.400	Moving Average: 0.238
Episodes 0750-0760	Max Reward: 0.400	Moving Average: 0.234
Episodes 0760-0770	Max Reward: 0.400	Moving Average: 0.233
Episodes 0770-0780	Max Reward: 0.400	Moving Average: 0.232
Episodes 0780-0790	Max Reward: 0.400	Moving Average: 0.235
Episodes 0790-0800	Max Reward: 0.400	Moving Average: 0.233
Episodes 0790-0800 Saved!!!	Max Reward: 0.400	Moving Average: 0.233
Episodes 0800-0810	Max Reward: 1.500	Moving Average: 0.249
Episodes 0810-0820	Max Reward: 0.900	Moving Average: 0.254
Episodes 0820-0830	Max Reward: 1.100	Moving Average: 0.273
Episodes 0830-0840	Max Reward: 1.800	Moving Average: 0.304
Episodes 0840-0850	Max Reward: 2.900	Moving Average: 0.380
Episodes 0850-0860	Max Reward: 1.200	Moving Average: 0.437
Episodes 0860-0870	Max Reward: 2.300	Moving Average: 0.461
Episodes 0870-0880	Max Reward: 3.000	Moving Average: 0.510
Episodes 0880-0890	Max Reward: 2.700	Moving Average: 0.542
Episodes 0890-0900	Max Reward: 5.200	Moving Average: 0.746

Episodes 0890-0900 Saved!!!	Max Reward: 5.200	Moving Average: 0.746
--------------------------------	-------------------	-----------------------

Episodes 0900-0910	Max Reward: 5.200	Moving Average: 0.846
Episodes 0910-0920	Max Reward: 4.700	Moving Average: 0.905
Episodes 0920-0930	Max Reward: 5.300	Moving Average: 1.115
Episodes 0930-0940	Max Reward: 2.900	Moving Average: 1.162
Episodes 0940-0950	Max Reward: 1.700	Moving Average: 1.122
Episodes 0950-0960	Max Reward: 3.000	Moving Average: 1.147
Episodes 0960-0970	Max Reward: 2.100	Moving Average: 1.218
Episodes 0970-0980	Max Reward: 2.100	Moving Average: 1.219
Episodes 0980-0990	Max Reward: 4.600	Moving Average: 1.257
Episodes 0990-1000	Max Reward: 3.400	Moving Average: 1.144

Episodes 0990-1000 Saved!!!	Max Reward: 3.400	Moving Average: 1.144
--------------------------------	-------------------	-----------------------

Episodes 1000-1010	Max Reward: 5.200	Moving Average: 1.088
Episodes 1010-1020	Max Reward: 5.200	Moving Average: 1.194
Episodes 1020-1030	Max Reward: 4.100	Moving Average: 1.077
Episodes 1030-1040	Max Reward: 2.300	Moving Average: 1.049
Episodes 1040-1050	Max Reward: 5.100	Moving Average: 1.110
Episodes 1050-1060	Max Reward: 4.100	Moving Average: 1.147
Episodes 1060-1070	Max Reward: 4.700	Moving Average: 1.203
Episodes 1070-1080	Max Reward: 2.000	Moving Average: 1.208
Episodes 1080-1090	Max Reward: 3.700	Moving Average: 1.276
Episodes 1090-1100	Max Reward: 5.200	Moving Average: 1.488

Episodes 1090-1100 Saved!!!	Max Reward: 5.200	Moving Average: 1.488
--------------------------------	-------------------	-----------------------

Episodes 1100-1110	Max Reward: 5.200	Moving Average: 1.659
Episodes 1110-1120	Max Reward: 5.200	Moving Average: 1.800
Episodes 1120-1130	Max Reward: 5.200	Moving Average: 1.947
Episodes 1130-1140	Max Reward: 5.200	Moving Average: 2.030
Episodes 1140-1150	Max Reward: 5.200	Moving Average: 2.094
Episodes 1150-1160	Max Reward: 5.200	Moving Average: 2.254
Episodes 1160-1170	Max Reward: 5.200	Moving Average: 2.462

<-- Environment solved in 1071 episodes!

<-- Moving Average: 2.505 over past 100 episodes

In [13]:


```

# plot the scores
from collections import deque
import matplotlib.pyplot as plt
import numpy as np
import random
import time
import torch
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='MADDPG')
plt.plot(np.arange(len(scores)), avgs, c='r', label='moving avg')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left');
fig.show()
fig.savefig("MADDPG_1e04__3000.png")

```

/opt/conda/lib/python3.6/site-packages/matplotlib/figure.py:418: UserWarning: matplotlib is currently using a non-GUI backend, so cannot show the figure

"matplotlib is currently using a non-GUI backend, "



In [12]:

```

## reinitialize the agents (if needed)
agent_0 = Agent(state_size, action_size, num_agents=1, random_seed=0)
agent_1 = Agent(state_size, action_size, num_agents=1, random_seed=0)

# Load the weights from file
agent_0_weights = 'checkpoint_actor_0.pth'
agent_1_weights = 'checkpoint_actor_1.pth'
agent_0.actor_local.load_state_dict(torch.load(agent_0_weights))
agent_1.actor_local.load_state_dict(torch.load(agent_1_weights))

```


In [14]:

```

CONSEC_EPISODES = 50
PRINT_EVERY = 1
ADD_NOISE = False

def test(n_episodes=100, max_t=1000, train_mode=False):

    scores_window = deque(maxlen=CONSEC_EPISODES)
    scores_all = []
    moving_average = []

    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=train_mode)[brain_name]          # reset the env
        states = np.reshape(env_info.vector_observations, (1,48)) # get states and combine them
        scores = np.zeros(num_agents)
        while True:
            actions = get_actions(states, ADD_NOISE)          # choose agent actions and combine them
            env_info = env.step(actions)[brain_name]          # send both agents' actions together to the environment
            next_states = np.reshape(env_info.vector_observations, (1, 48)) # combine the agent next states
            rewards = env_info.rewards          # get reward
            done = env_info.local_done          # see if episode finished
            scores += np.max(rewards)          # update the score for each agent
            states = next_states          # roll over states to next time step
            if np.any(done):          # exit loop if episode finished
                break

            ep_best_score = np.max(scores)
            scores_window.append(ep_best_score)
            scores_all.append(ep_best_score)
            moving_average.append(np.mean(scores_window))

        # print results
        if i_episode % PRINT_EVERY == 0:
            print('Episodes {:0>4d}-{:0>4d}\tMax Reward: {:.3f}\tMoving Average: {:.3f}'.format(
                i_episode-PRINT_EVERY, i_episode, np.max(scores_all[-PRINT_EVERY:]), moving_average[-1]))

    return scores_all, moving_average

```

In [15]:

```
scores, avgs = test()

# plot the scores
from collections import deque
import matplotlib.pyplot as plt
import numpy as np
import random
import time
import torch

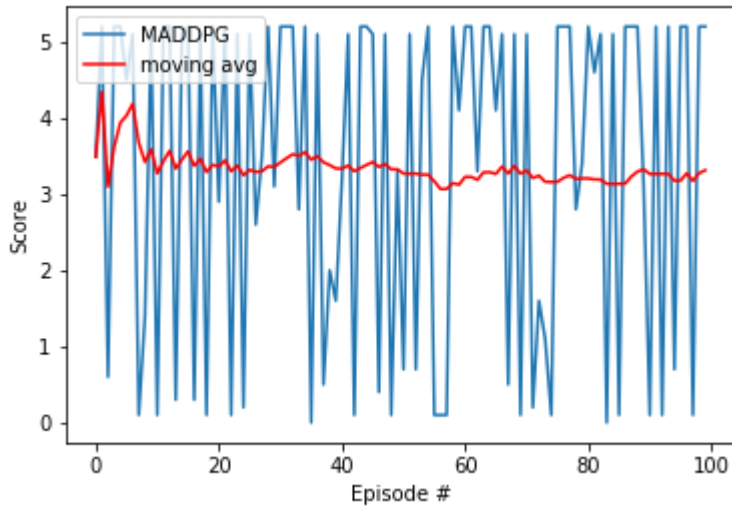
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='MADDPG')
plt.plot(np.arange(len(scores)), avgs, c='r', label='moving avg')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left');
fig.savefig("MADDPG_Test2.png")
fig.show()
```

Episodes 0000-0001	Max Reward: 3.490	Moving Average: 3.490
Episodes 0001-0002	Max Reward: 5.200	Moving Average: 4.345
Episodes 0002-0003	Max Reward: 0.600	Moving Average: 3.097
Episodes 0003-0004	Max Reward: 5.200	Moving Average: 3.623
Episodes 0004-0005	Max Reward: 5.200	Moving Average: 3.938
Episodes 0005-0006	Max Reward: 4.500	Moving Average: 4.032
Episodes 0006-0007	Max Reward: 5.100	Moving Average: 4.184
Episodes 0007-0008	Max Reward: 0.100	Moving Average: 3.674
Episodes 0008-0009	Max Reward: 1.400	Moving Average: 3.421
Episodes 0009-0010	Max Reward: 5.100	Moving Average: 3.589
Episodes 0010-0011	Max Reward: 0.100	Moving Average: 3.272
Episodes 0011-0012	Max Reward: 5.200	Moving Average: 3.433
Episodes 0012-0013	Max Reward: 5.200	Moving Average: 3.568
Episodes 0013-0014	Max Reward: 0.300	Moving Average: 3.335
Episodes 0014-0015	Max Reward: 5.100	Moving Average: 3.453
Episodes 0015-0016	Max Reward: 5.200	Moving Average: 3.562
Episodes 0016-0017	Max Reward: 0.300	Moving Average: 3.370
Episodes 0017-0018	Max Reward: 5.100	Moving Average: 3.466
Episodes 0018-0019	Max Reward: 0.100	Moving Average: 3.289
Episodes 0019-0020	Max Reward: 5.200	Moving Average: 3.385
Episodes 0020-0021	Max Reward: 2.900	Moving Average: 3.361
Episodes 0021-0022	Max Reward: 5.200	Moving Average: 3.445
Episodes 0022-0023	Max Reward: 0.100	Moving Average: 3.300
Episodes 0023-0024	Max Reward: 5.100	Moving Average: 3.375
Episodes 0024-0025	Max Reward: 0.200	Moving Average: 3.248
Episodes 0025-0026	Max Reward: 5.100	Moving Average: 3.319
Episodes 0026-0027	Max Reward: 2.600	Moving Average: 3.292
Episodes 0027-0028	Max Reward: 3.500	Moving Average: 3.300
Episodes 0028-0029	Max Reward: 5.200	Moving Average: 3.365
Episodes 0029-0030	Max Reward: 3.100	Moving Average: 3.356
Episodes 0030-0031	Max Reward: 5.200	Moving Average: 3.416
Episodes 0031-0032	Max Reward: 5.200	Moving Average: 3.472
Episodes 0032-0033	Max Reward: 5.200	Moving Average: 3.524
Episodes 0033-0034	Max Reward: 2.800	Moving Average: 3.503
Episodes 0034-0035	Max Reward: 5.200	Moving Average: 3.551
Episodes 0035-0036	Max Reward: 0.000	Moving Average: 3.453
Episodes 0036-0037	Max Reward: 5.100	Moving Average: 3.497
Episodes 0037-0038	Max Reward: 0.500	Moving Average: 3.418
Episodes 0038-0039	Max Reward: 2.000	Moving Average: 3.382
Episodes 0039-0040	Max Reward: 1.600	Moving Average: 3.337
Episodes 0040-0041	Max Reward: 3.200	Moving Average: 3.334
Episodes 0041-0042	Max Reward: 5.100	Moving Average: 3.376
Episodes 0042-0043	Max Reward: 0.100	Moving Average: 3.300
Episodes 0043-0044	Max Reward: 5.200	Moving Average: 3.343
Episodes 0044-0045	Max Reward: 5.200	Moving Average: 3.384
Episodes 0045-0046	Max Reward: 5.100	Moving Average: 3.422
Episodes 0046-0047	Max Reward: 0.400	Moving Average: 3.357
Episodes 0047-0048	Max Reward: 5.100	Moving Average: 3.394
Episodes 0048-0049	Max Reward: 0.100	Moving Average: 3.326
Episodes 0049-0050	Max Reward: 3.300	Moving Average: 3.326
Episodes 0050-0051	Max Reward: 0.700	Moving Average: 3.270
Episodes 0051-0052	Max Reward: 5.100	Moving Average: 3.268
Episodes 0052-0053	Max Reward: 0.700	Moving Average: 3.270
Episodes 0053-0054	Max Reward: 4.500	Moving Average: 3.256
Episodes 0054-0055	Max Reward: 5.200	Moving Average: 3.256
Episodes 0055-0056	Max Reward: 0.100	Moving Average: 3.168
Episodes 0056-0057	Max Reward: 0.100	Moving Average: 3.068
Episodes 0057-0058	Max Reward: 0.100	Moving Average: 3.068
Episodes 0058-0059	Max Reward: 5.200	Moving Average: 3.144
Episodes 0059-0060	Max Reward: 4.100	Moving Average: 3.124
Episodes 0060-0061	Max Reward: 5.200	Moving Average: 3.226

Episodes 0061-0062	Max Reward: 5.200	Moving Average: 3.226
Episodes 0062-0063	Max Reward: 3.300	Moving Average: 3.188
Episodes 0063-0064	Max Reward: 5.200	Moving Average: 3.286
Episodes 0064-0065	Max Reward: 5.200	Moving Average: 3.288
Episodes 0065-0066	Max Reward: 4.100	Moving Average: 3.266
Episodes 0066-0067	Max Reward: 5.100	Moving Average: 3.362
Episodes 0067-0068	Max Reward: 0.500	Moving Average: 3.270
Episodes 0068-0069	Max Reward: 5.100	Moving Average: 3.370
Episodes 0069-0070	Max Reward: 0.100	Moving Average: 3.268
Episodes 0070-0071	Max Reward: 5.100	Moving Average: 3.312
Episodes 0071-0072	Max Reward: 0.200	Moving Average: 3.212
Episodes 0072-0073	Max Reward: 1.600	Moving Average: 3.242
Episodes 0073-0074	Max Reward: 1.100	Moving Average: 3.162
Episodes 0074-0075	Max Reward: 0.100	Moving Average: 3.160
Episodes 0075-0076	Max Reward: 5.200	Moving Average: 3.162
Episodes 0076-0077	Max Reward: 5.200	Moving Average: 3.214
Episodes 0077-0078	Max Reward: 5.200	Moving Average: 3.248
Episodes 0078-0079	Max Reward: 2.800	Moving Average: 3.200
Episodes 0079-0080	Max Reward: 3.400	Moving Average: 3.206
Episodes 0080-0081	Max Reward: 5.200	Moving Average: 3.206
Episodes 0081-0082	Max Reward: 4.600	Moving Average: 3.194
Episodes 0082-0083	Max Reward: 5.100	Moving Average: 3.192
Episodes 0083-0084	Max Reward: 0.000	Moving Average: 3.136
Episodes 0084-0085	Max Reward: 5.100	Moving Average: 3.134
Episodes 0085-0086	Max Reward: 0.100	Moving Average: 3.136
Episodes 0086-0087	Max Reward: 5.200	Moving Average: 3.138
Episodes 0087-0088	Max Reward: 5.200	Moving Average: 3.232
Episodes 0088-0089	Max Reward: 5.200	Moving Average: 3.296
Episodes 0089-0090	Max Reward: 3.100	Moving Average: 3.326
Episodes 0090-0091	Max Reward: 0.100	Moving Average: 3.264
Episodes 0091-0092	Max Reward: 5.200	Moving Average: 3.266
Episodes 0092-0093	Max Reward: 0.100	Moving Average: 3.266
Episodes 0093-0094	Max Reward: 5.200	Moving Average: 3.266
Episodes 0094-0095	Max Reward: 0.700	Moving Average: 3.176
Episodes 0095-0096	Max Reward: 5.200	Moving Average: 3.178
Episodes 0096-0097	Max Reward: 5.200	Moving Average: 3.274
Episodes 0097-0098	Max Reward: 0.100	Moving Average: 3.174
Episodes 0098-0099	Max Reward: 5.200	Moving Average: 3.276
Episodes 0099-0100	Max Reward: 5.200	Moving Average: 3.314

/opt/conda/lib/python3.6/site-packages/matplotlib/figure.py:418: UserWarning: matplotlib is currently using a non-GUI backend, so cannot show the figure

"matplotlib is currently using a non-GUI backend, "

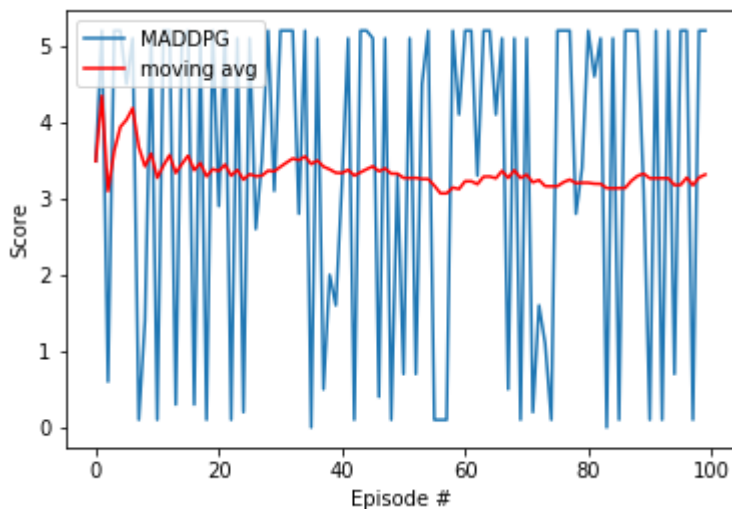


In [16]:

```
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='MADDPG')
plt.plot(np.arange(len(scores)), avgs, c='r', label='moving avg')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left')
fig.savefig("MADDPG_Test2.png")
fig.show()
```

/opt/conda/lib/python3.6/site-packages/matplotlib/figure.py:418: UserWarning: matplotlib is currently using a non-GUI backend, so cannot show the figure

"matplotlib is currently using a non-GUI backend, "



In [17]:

```
#PERFORMANCE IMPROVEMENT  
# LEARNING RATE WITH 1E-5 OR LESS  
# KEEPING LEARNING RATE DIFFERENT BETWEEN 2 AGENTS  
# INCREASE EPOCHS  
# DENSE NETWORK
```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: